

Parallel Computer Architecture

A Hardware / Software Approach

by Culler and Singh, Morgan Kaufmann/Elsevier, 1998

Chapter2 Parallel Programs



The University of
British Columbia

Overview

Last lecture set...

- Optimizing your sequential program
 - Compiler optimizations
 - Measuring performance of a single program
 - Summarizing performance – **distill** set of performance results into 1 magic number

This set...

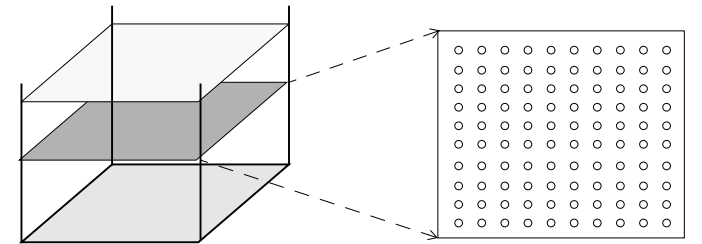
- Parallel programming
 - A general introduction
 - Readings – Culler text, Chapters 2 & 3

Parallel Programming

Culler, Chapter 2

Case Studies

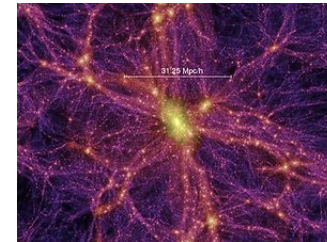
- Simulating Ocean Currents
(Regular grid)



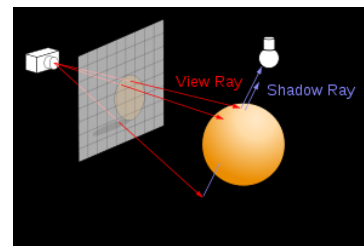
(a) Cross-sections

(b) Spatial discretization of a cross-section

- Galaxy Simulation
(Irregular)



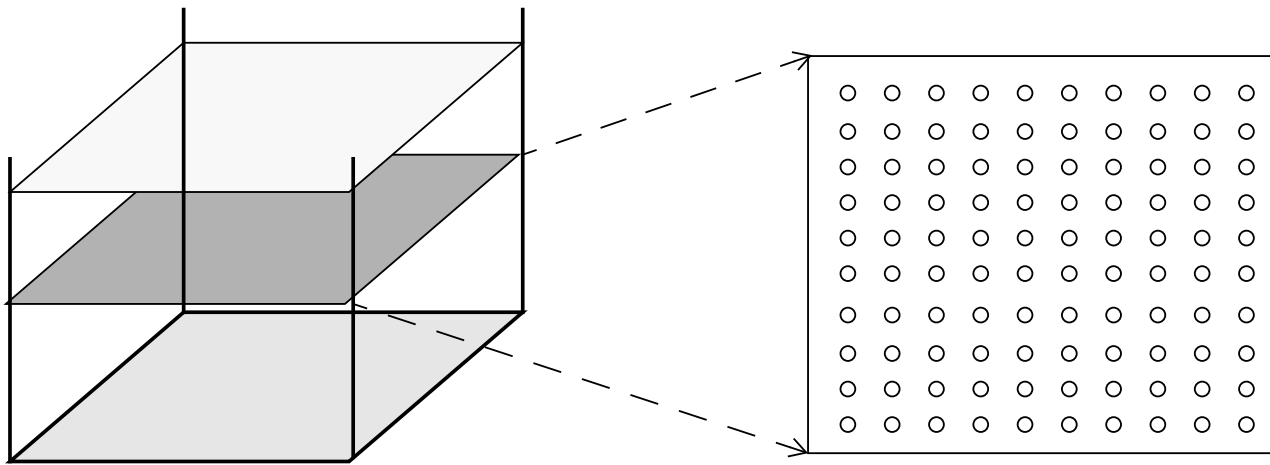
- Ray Tracing
(individual beams)



- Data Mining
(dependencies among tables)

Parallelizable Applications

- Ocean (similar to weather prediction)
 - Model current in ocean
 - Multi-dimensional grid points at each time stamp

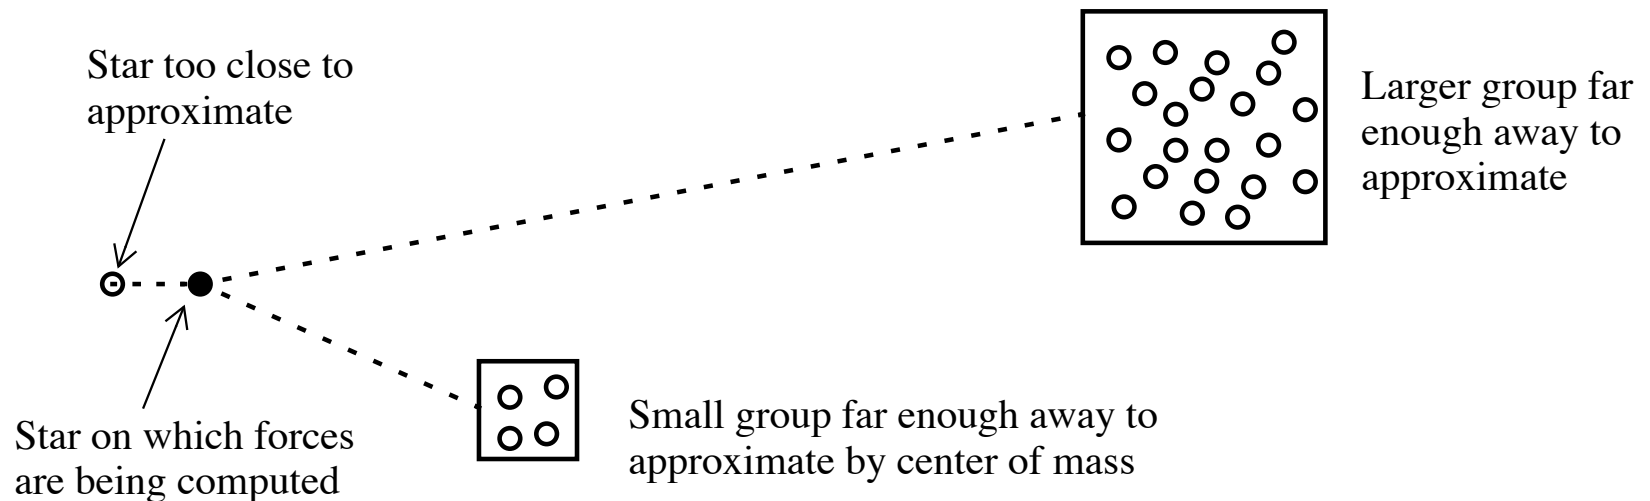


(a) Cross-sections

(b) Spatial discretization of a cross-section

Parallelizable Applications

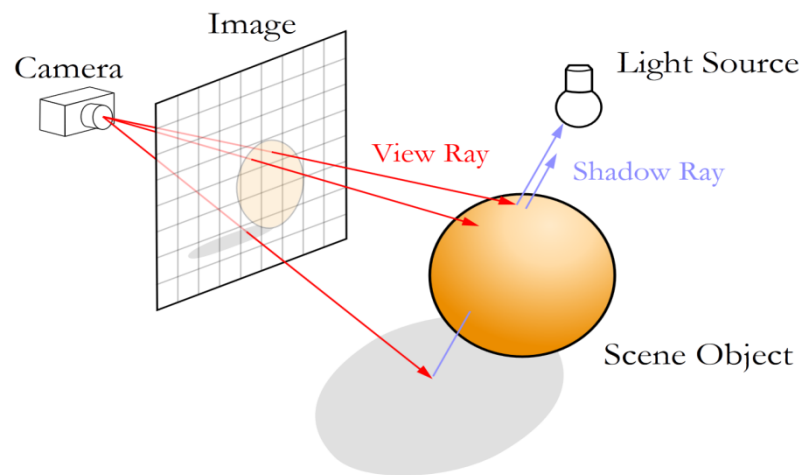
- Galaxies (N-body simulation)
 - Model force interactions among stars
 - Barnes-Hut:
 - Speed-up through approximation: group nearby stars together



Raytrace Simulation

Rendering scenes by tracing:

- Produces high quality images but very expensive in computations
- Represent image as two-dimensional array of pixels
- Variables calculated for each pixel: color, opacity and brightness
- Shoot rays into scene through image plane, they strike and bounce
- Compute reflection, refraction and lighting interactions
- Concurrency is found for rays shot through different pixels



Source: Wikipedia

Parallelizable Applications

- Data-Mining
 - Analyze correlations of attributes in a database
 - What is probability of:
 - buying item T2
 - given customer has bought item T1 ?
 - Determine itemset:
 - Set of k elements purchased together
 - Given that customer:
 - has set of items S1 in cart
 - what is probability of buying items in set S2 ?

Parallelization Process

- At a high level, parallelization involves
 - Identifying **what work can be done in parallel**,
 - Determining **how to distribute the work and** perhaps the **data** among the processing nodes, and
 - **Managing** the necessary data **access, communication** and **synchronization**
- The of parallelization may be performed
 - Programmer
 - Compiler
 - Runtime system
 - Operating system

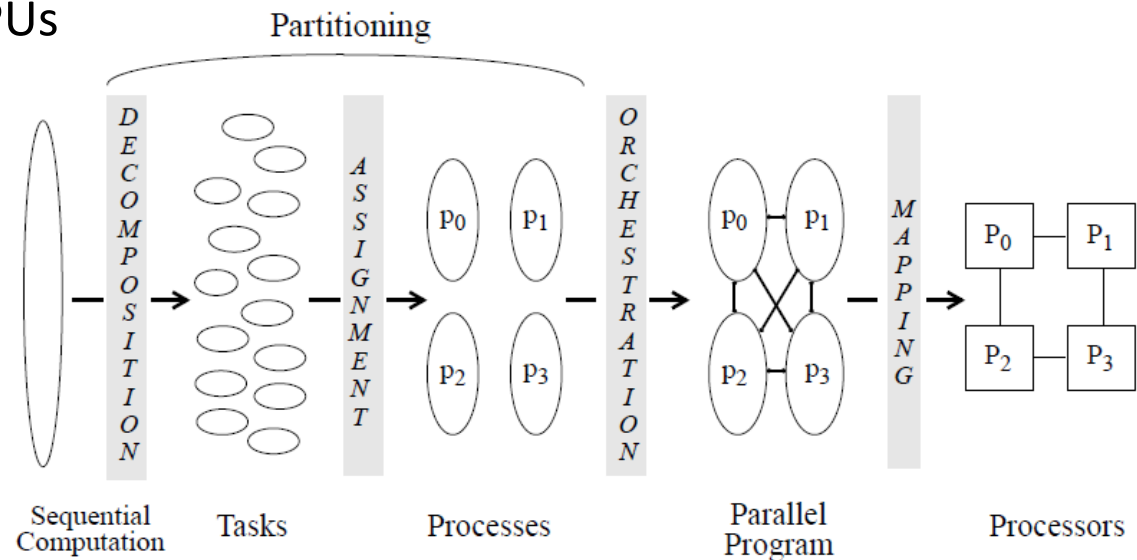
Parallelization Process

- **Important concepts**
 - Task
(Smallest unit of concurrency)
 - Process
(An abstract entity that performs a subset of tasks)
 - Processor
(Physical resource executing processes)

Parallelization Process

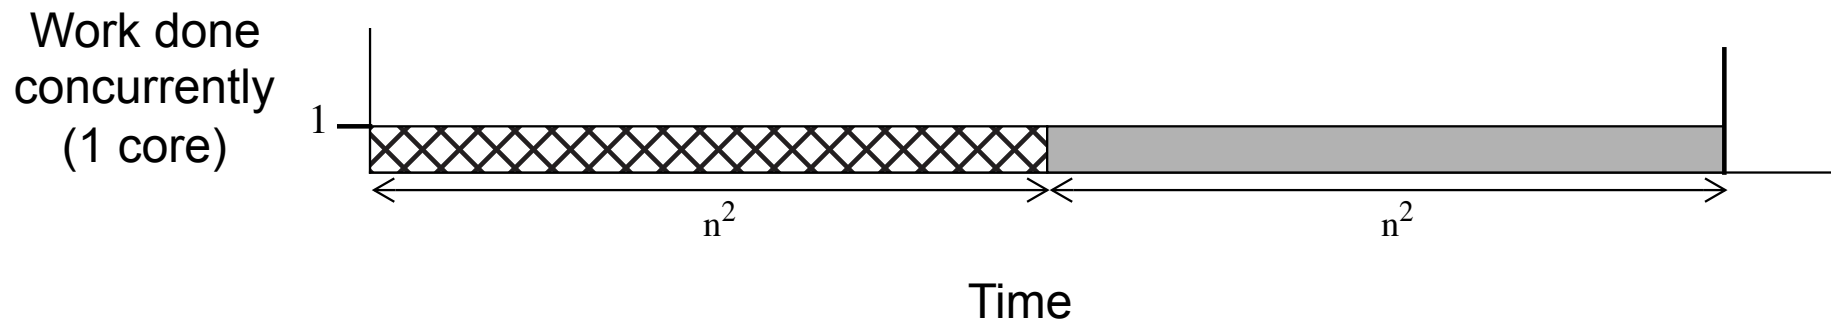
Steps of Parallelization

- Decomposition (into tasks)
- Assignment (group of tasks to processes)
- Orchestration (arrange communication between processes)
 - Select programming model: message passing and/or shared memory
- Mapping (assign processes to physical resources)
 - Memory and CPUs

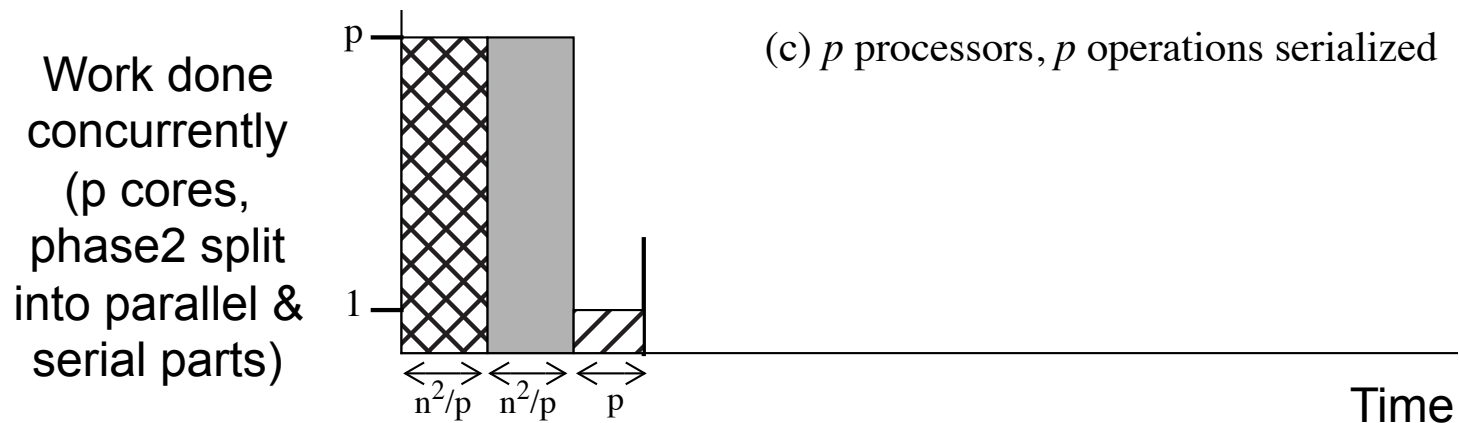
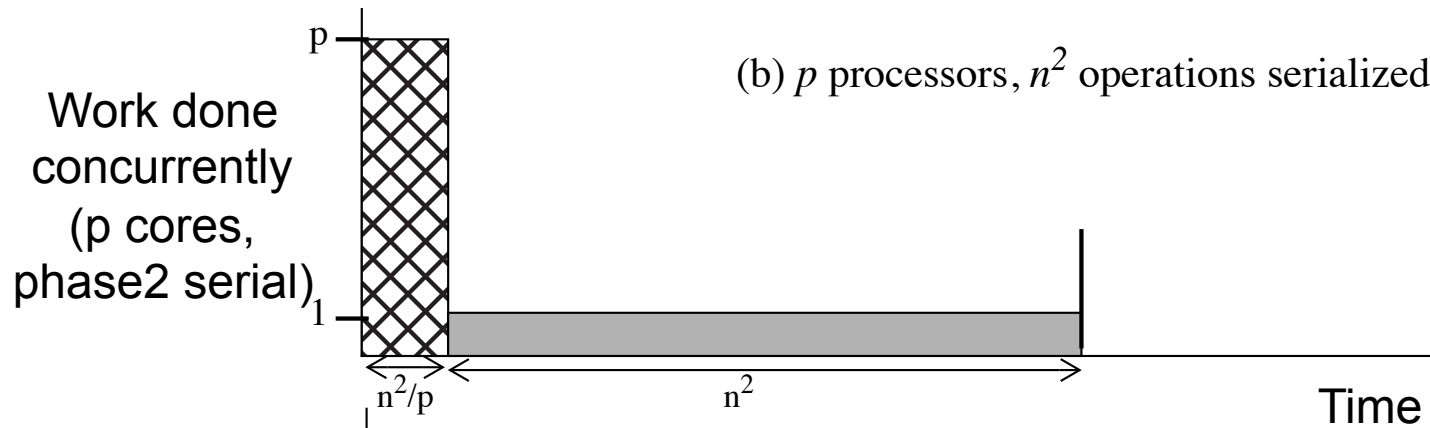


Parallelization and Limits

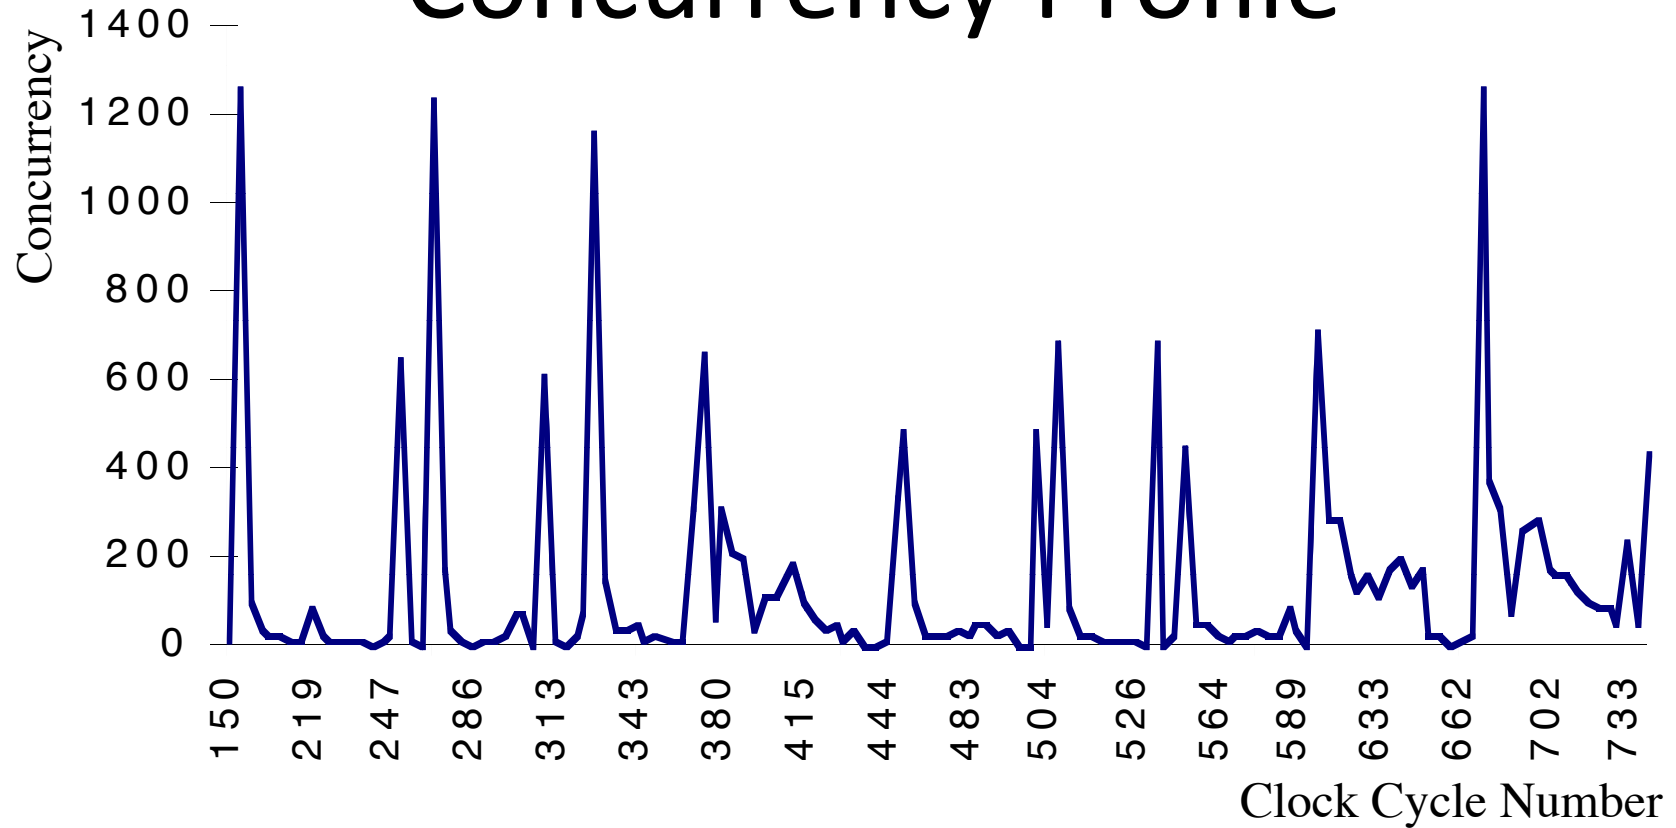
- Example Problem
 - Two-phase algorithm
 - Phase 1: independently update n^2 data points
 - Phase 2: sum all n^2 data points



Parallelization and Limits

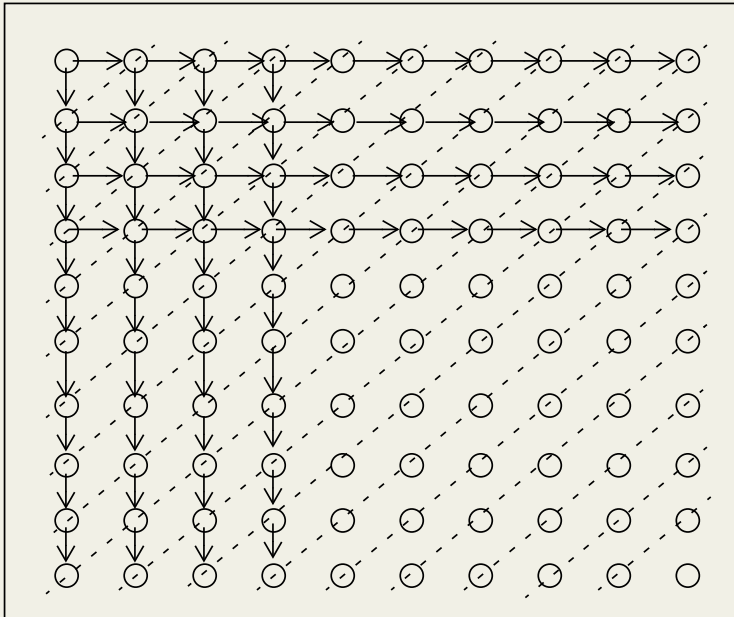
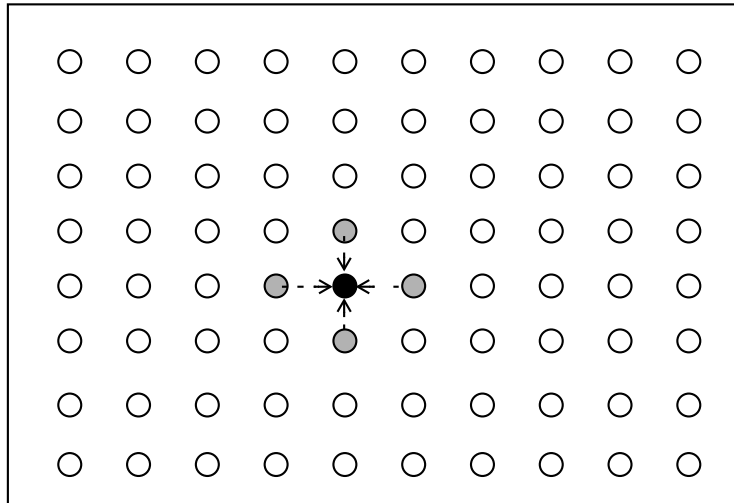


Concurrency Profile



Concurrency profile for logic simulator
(Fig 2.5 in Culler)

Case Study: Ocean



Do

err=0;

For each data point in grid

temp = A[];

A[] = 0.2 * (... A[] ...);

err += abs(A[] - temp);

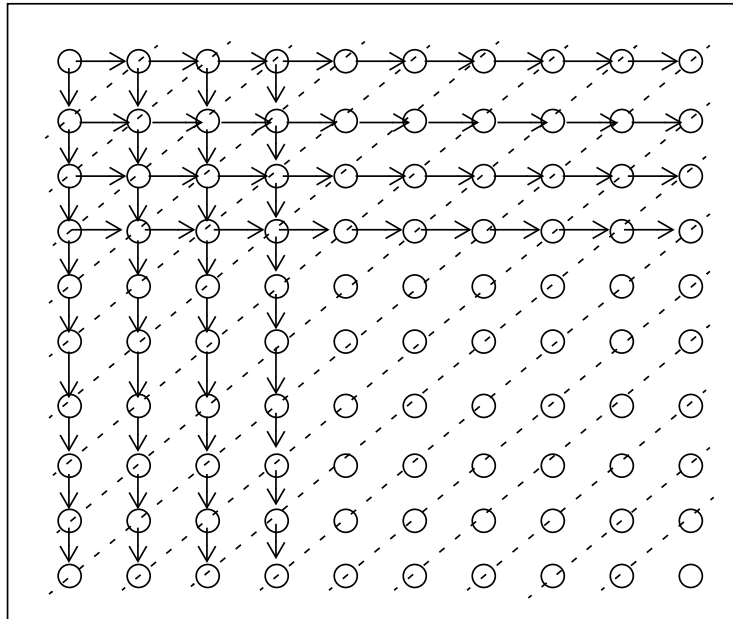
End for

Until err < tolerance

In the sequential algorithm, we execute row-by-row, top-down

Ocean: Decomposition

- Find the parallelism
- Can we divide grid points into groups
 - Compute each group on different processors?



Dependences and concurrency in the Gauss-Seidel equation solver computation.

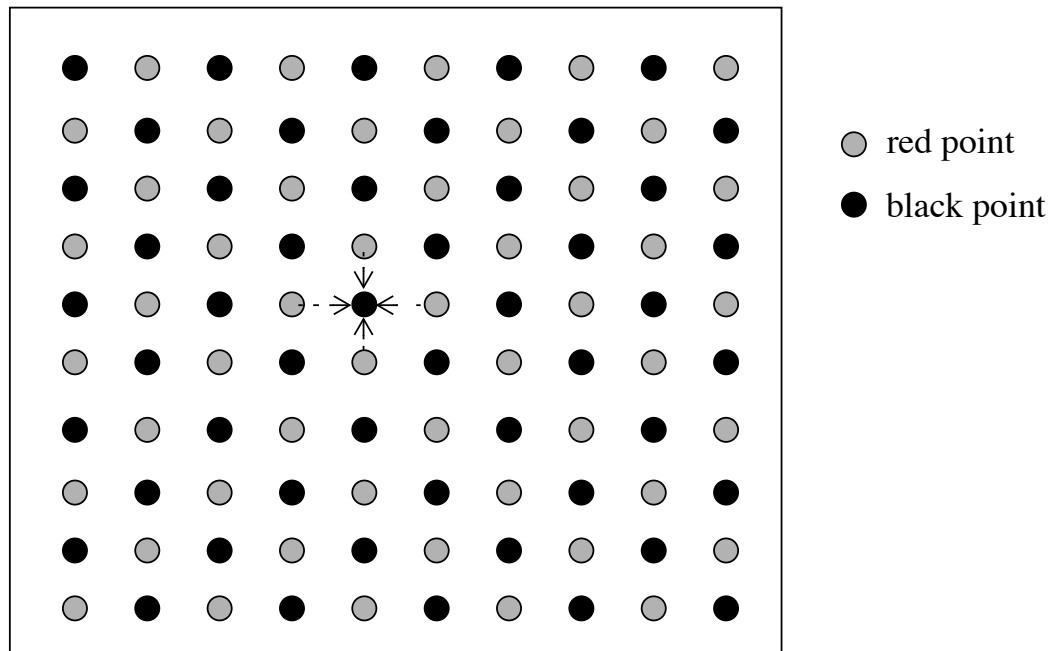
- Obstacles:
 - Each point has a data dependency on its neighbors
 - The final summation of data at each point

Solution 1

- Analyze the dependence pattern
- Identify: no dependence along diagonal line
 - Each diagonal computed in sequence
 - Within each diagonal, points can be processed in parallel
- Problems
 - Load imbalance
 - Diagonal changes length each iteration !
 - Average parallelism $N/2$
 - Synchronization overhead

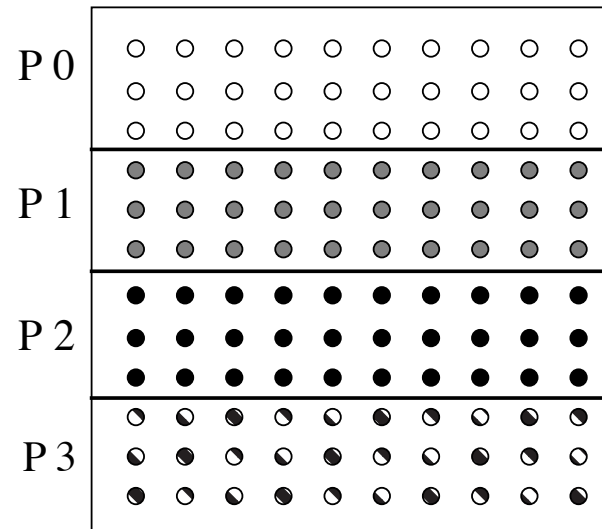
Solution 2

- Divide grid points to dark and light colors in alternating order
 - Points of the same color are independent of each other
 - Average parallelism is $N^2/2$
- NB: This is a different algorithm, produces a different result!
 - Computation may be different, but still solves the problem (!)
 - More parallelism, but requires more iterations to converge



Solution 3

- Each process assigned contiguous, equal number of rows

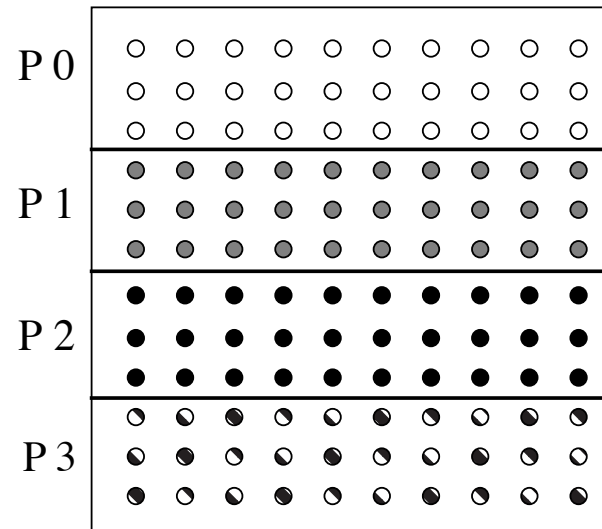


- Note: this will not deliver the same result that we get from the sequential implementation!
- Communication required for adjacent rows
 - May need “ghost rows”: a copy of adjacent rows to use locally
- Let’s use this version going forward...

Shared Memory

Solution 3

- Each process assigned contiguous, equal number of rows



- Communication required for rows at the partition boundaries**

Ocean: Orchestration in Shared Memory

- Shared memory
 - Each process can access all data (any partition)
- *Barrier* is used to ensure all processes proceed in lock-step to the next sweep (iteration #)
- Requires *mutual exclusion* to modify any shared (global) variables
 - Eg, the final summation

Shared Memory

Parallel Programming Primitives

Table 2-2 Key Shared Address Space Primitives

Name	Syntax	Function
CREATE	<code>CREATE(p,proc,args)</code>	Create <code>p</code> processes that start executing at procedure <code>proc</code> with arguments <code>args</code> .
G_MALLOC	<code>G_MALLOC(size)</code>	Allocate shared data of <code>size</code> bytes
LOCK	<code>LOCK(name)</code>	Acquire mutually exclusive access
UNLOCK	<code>UNLOCK(name)</code>	Release mutually exclusive access
BARRIER	<code>BARRIER(name, number)</code>	Global synchronization among <code>number</code> processes: None gets past BARRIER until <code>number</code> have arrived
WAIT_FOR_END	<code>WAIT_FOR_END(number)</code>	Wait for <code>number</code> processes to terminate
wait for flag	<code>while (!flag);</code> or <code>WAIT(flag)</code>	Wait for <code>flag</code> to be set (spin or block); for point-to-point event synchronization.
set flag	<code>flag = 1;</code> or <code>SIGNAL(flag)</code>	Set <code>flag</code> ; wakes up process spinning or blocked on <code>flag</code> , if any

Orchestration – Shared memory

```

1. int n, nprocs;           /* matrix dimension and number of processors to be used */
2a. float **A, diff;        /* A is global (shared) array representing the grid */
                             /* diff is global (shared) maximum difference in current sweep */
2b. LOCKDEC(diff_lock);     /* declaration of lock to enforce mutual exclusion */
2c. BARDEC (bar1);          /* barrier declaration for global synchronization between sweeps */

3. main()
4. begin
5.   read(n);   read(nprocs);   /* read input matrix size and number of processes */
6.   A ← G_MALLOC (a two-dimensional array of size n+2 by n+2 doubles);
7.   initialize(A);             /* initialize A in an unspecified way */
8a.  CREATE (nprocs-1, Solve, A);
8   Solve(A);                  /* main process becomes a worker too */
8b.  WAIT_FOR_END;             /* wait for all child processes created to terminate */
9. end main

10. procedure Solve(A)
11. float **A;                 /* A is entire n+2-by-n+2 shared array, as in the sequential program */
12. begin
13.  int i,j, pid, done = 0;
14.  float temp, mydiff = 0;    /* private variables */
14a. int mymin ← 1 + (pid * n/nprocs); /* assume that n is exactly divisible by */
14b. int mymax ← mymin + n/nprocs - 1; /* nprocs for simplicity here */

15. while (!done) do           /* outer loop over all diagonal elements */
16.   mydiff = diff = 0;        /* set global diff to 0 (okay for all to do it) */
17.   for i ← mymin to mymax do /* for each of my rows */
18.     for j ← 1 to n do       /* for all elements in that row */
19.       temp = A[i,j];
20.       A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
21.         A[i,j+1] + A[i+1,j]);
22.       mydiff += abs(A[i,j] - temp);
23.     endfor
24.   endfor

25a. LOCK(diff_lock);          /* update global diff if necessary */
25b.   diff += mydiff;
25c. UNLOCK(diff_lock);
25d. BARRIER(bar1, nprocs);   /* ensure all have got here before checking if done */

25e. if (diff/(n*n) < TOL) then done = 1; /* check convergence; all get same answer */
25f. BARRIER(bar1, nprocs);     /* see Exercise c */
26. endwhile
27. end procedure

```


Orchestration – Shared memory

```
1. int n, nprocs;           /* matrix dimension and number of processors to be used */
2a. float **A, diff;        /* A is global (shared) array representing the grid */
                             /* diff is global (shared) maximum difference in current sweep */
2b. LOCKDEC(diff_lock);     /* declaration of lock to enforce mutual exclusion */
2c. BARDEC (bar1);          /* barrier declaration for global synchronization between sweeps */
```

What is a lock? (mutual exclusion)

Only one process can "acquire" lock at a time
Other processes wait until lock is available

What is a barrier? (wait for rendezvous)

Each process stops & waits
After all processes reach barrier, they can proceed

```
3. main()
4. begin
5.   read(n);   read(nprocs);   /* read input matrix size and number of processes */
6.   A ← G_MALLOC (a two-dimensional array of size n+2 by n+2 doubles);
7.   initialize(A);             /* initialize A in an unspecified way */
```

G_MALLOC: allocate from shared memory region

Orchestration – Shared memory

```
8a.  CREATE (nprocs-1, Solve, A);  
8    Solve(A);                      /* main process becomes a worker too */  
8b.  WAIT_FOR_END;                  /* wait for all child processes created to terminate */  
9.  end main
```

- CREATE(p,proc,args)
 - Create p processes
- WAIT_FOR_END(number)
 - Wait for number processes to terminate
- Pseudo code is implementable in Pthread library in C:

CREATE

```
int pthread_create(pthread_t *restrict thread, const  
pthread_attr_t *restrict attr, void *(*start_routine)(void *),  
void *restrict arg);
```

WAIT_FOR_END

```
int pthread_join(pthread_t thread, void **value_ptr);
```

Orchestration – Shared memory

```
15. while (!done) do                                /* outer loop over all diagonal elements */
16.   mydiff = diff = 0;                             /* set global diff to 0 (okay for all to do it) */
17.   for i ← mymin to mymax do                       /* for each of my rows */
18.     for j ← 1 to n do                             /* for all elements in that row */
19.       temp = A[i,j];
20.       A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
21.         A[i,j+1] + A[i+1,j]);
22.       mydiff += abs(A[i,j] - temp);
23.     endfor
24.   endfor
```

Main algorithm

- Similar to serial code
- Each thread computes one part of the array
- Each thread computes the diff of its part (local variable “mydiff”)

Orchestration – Shared memory

- Serial part – how to add all mydiffs – needs a LOCK

```
25a. LOCK(diff_lock);           /* update global diff if necessary */
25b.     diff += mydiff;
25c. UNLOCK(diff_lock);
```

- Why do we need a lock?

<u>P1</u>	<u>P2</u>	
$r1 \leftarrow \text{diff}$		<i>{P1 gets 0 in its r1}</i>
	$r1 \leftarrow \text{diff}$	<i>{P2 also gets 0}</i>
$r1 \leftarrow r1 + r2$		<i>{P1 sets its r1 to 1}</i>
	$r1 \leftarrow r1 + r2$	<i>{P2 sets its r1 to 1}</i>
$\text{diff} \leftarrow r1$		<i>{P1 sets cell_cost to 1}</i>
	$\text{diff} \leftarrow r1$	<i>{P2 also sets cell_cost to 1}</i>

- each CPU needs to execute this read/add/write atomically

Orchestration – Shared memory

Process/Thread Synchronization:

```
25d. BARRIER(bar1, nprocs);           /* ensure all have got here before checking if done*/  
  
25e. if (diff/(n*n) < TOL) then done = 1; /* check convergence; all get same answer*/  
25f. BARRIER(bar1, nprocs);           /* see Exercise c */
```

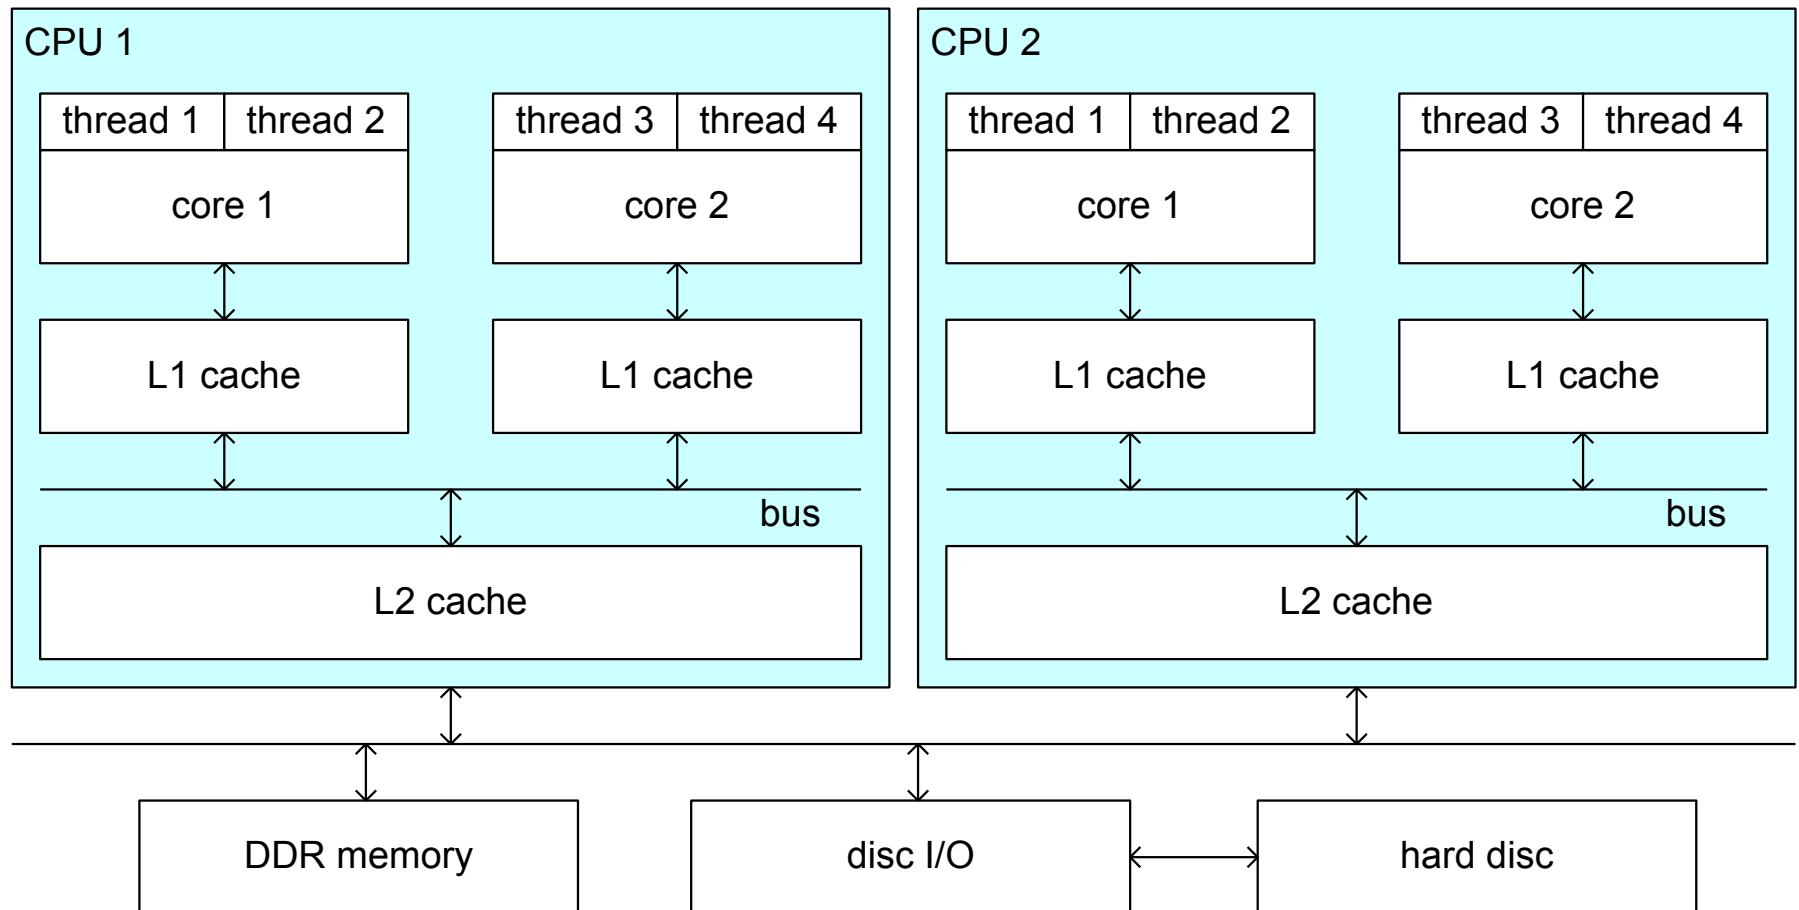
- First barrier
 - Ensures all threads update diff
- Second barrier
 - Get ready for next iteration

Main Part + Lock + Barriers

```
15. while (!done) do                                /* outer loop over all diagonal elements */
16.   mydiff = diff = 0;                             /* set global diff to 0 (okay for all to do it) */
17.   for i ← mymin to mymax do                       /* for each of my rows */
18.     for j ← 1 to n do                             /* for all elements in that row */
19.       temp = A[i,j];
20.       A[i,j] = 0.2 * (A[i,j] + A[i,j-1] + A[i-1,j] +
21.         A[i,j+1] + A[i+1,j]);
22.       mydiff += abs(A[i,j] - temp);
23.     endfor
24.   endfor
25a.  LOCK(diff_lock);                               /* update global diff if necessary */
25b.    diff += mydiff;
25c.  UNLOCK(diff_lock);
25d.  BARRIER(bar1, nprocs);                       /* ensure all have got here before checking if done */
25e.  if (diff/(n*n) < TOL) then done = 1; /* check convergence; all get same answer */
25f.  BARRIER(bar1, nprocs);                       /* see Exercise c */
26. endwhile
27. end procedure
```

Orchestration – Shared memory

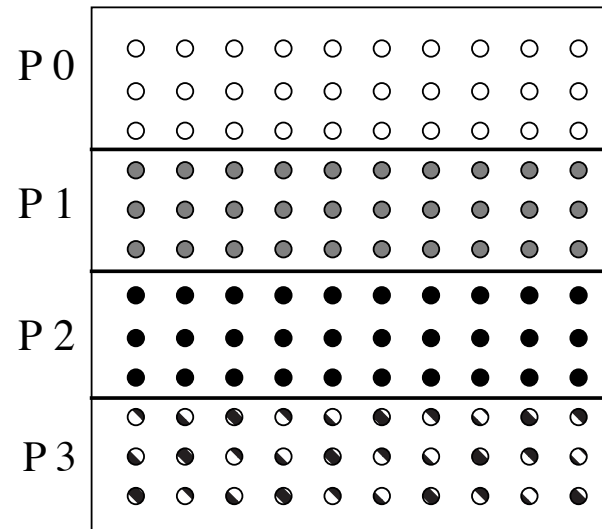
- Note that you have a memory hierarchy: data locality?
Commonly: Non-Uniform Memory Access (NUMA)



Message Passing

Solution 3

- Each process assigned contiguous, equal number of rows



- Communication required for rows at the partition boundaries
 - Must send messages containing “row updates” to neighbours

Ocean:

Orchestration in Message Passing

- Message passing
 - Each partition of points is stored on separate computer (private address space)
- No global variables / no shared memory
- Border rows need to be copied among neighbors
 - “Ghost rows”
- One process is responsible for global summation
 - Lock is implied, not explicit

Orchestration – Message passing

```
15. while (!done) do
16.     mydiff = 0;                                /*set local diff to 0*/
16a.    if (pid != 0) then SEND(&myA[1,0],n*sizeof(float),pid-1,ROW);
16b.    if (pid != nprocs-1) then
16b.1.        SEND(&myA[n',0],n*sizeof(float),pid+1,ROW);
16c.    if (pid != 0) then RECEIVE(&myA[0,0],n*sizeof(float),pid-1,ROW);
16d.    if (pid != nprocs-1) then
16d.1.        RECEIVE(&myA[n'+1,0],n*sizeof(float), pid+1,ROW);
16d.2.        /*border rows of neighbors have now been copied
16d.3.        into myA[0,*] and myA[n'+1,*]*/
17.    for i ← 1 to n' do                          /*for each of my (nonghost) rows*/
18.        for j ← 1 to n do                        /*for all nonborder elements in that row*/
19.            temp = myA[i,j];
20.            myA[i,j] = 0.2 * (myA[i,j] + myA[i,j-1] + myA[i-1,j] +
21.                myA[i,j+1] + myA[i+1,j]);
22.            mydiff += abs(myA[i,j] - temp);
23.        endfor
24.    endfor
25.    /*communicate local diff values and determine if
25.1.    done; can be replaced by reduction and broadcast*/
25a.    if (pid != 0) then                          /*process 0 holds global total diff*/
25b.        SEND(mydiff,sizeof(float),0,DIFF);
25c.        RECEIVE(done,sizeof(int),0,DONE);
25d.    else                                          /*pid 0 does this*/
25e.        for i ← 1 to nprocs-1 do                /*for each other process*/
25f.            RECEIVE(tempdiff,sizeof(float),*,DIFF);
25g.            mydiff += tempdiff;                  /*accumulate into total*/
25h.        endfor
25i.        if (mydiff/(n*n) < TOL) then done = 1;
25j.        for i ← 1 to nprocs-1 do                /*for each other process*/
25k.            SEND(done,sizeof(int),i,DONE);
25l.        endfor
25m.    endif
26. endwhile
27. end procedure
```

Message passing – update ghost rows

```
15. while (!done) do
16.     mydiff = 0;                                /*set local diff to 0*/
16a.    if (pid != 0) then SEND(&myA[1,0],n*sizeof(float),pid-1,ROW);
16b.    if (pid != nprocs-1) then
        SEND(&myA[n',0],n*sizeof(float),pid+1,ROW);
16c.    if (pid != 0) then RECEIVE(&myA[0,0],n*sizeof(float),pid-1,ROW);
16d.    if (pid != nprocs-1) then
        RECEIVE(&myA[n'+1,0],n*sizeof(float), pid+1,ROW);
        /*border rows of neighbors have now been copied
        into myA[0,*] and myA[n'+1,*]*/
```

Message passing – main algorithm

```
/*border rows of neighbors have now been copied
into myA[0,*] and myA[n'+1,*]*/
17.   for i ← 1 to n' do      /*for each of my (nonghost) rows*/
18.     for j ← 1 to n do    /*for all nonborder elements in that row*/
19.       temp = myA[i,j];
20.       myA[i,j] = 0.2 * (myA[i,j] + myA[i,j-1] + myA[i-1,j] +
21.         myA[i,j+1] + myA[i+1,j]);
22.       mydiff += abs(myA[i,j] - temp);
23.     endfor
24.   endfor

/*communicate local diff values and determine if
done; can be replaced by reduction and broadcast*/
```

Message passing – add “mydiff” errors

```
25a.  if (pid != 0) then                                /*process 0 holds global total diff*/
25b.      SEND(mydiff,sizeof(float),0,DIFF);
25c.      RECEIVE(done,sizeof(int),0,DONE);
25d.  else                                              /*pid 0 does this*/
25e.      for i ← 1 to nprocs-1 do                    /*for each other process*/
25f.          RECEIVE(tempdiff,sizeof(float),*,DIFF);
25g.          mydiff += tempdiff;                      /*accumulate into total*/
25h.      endfor
```

Orchestration – are we done yet?

```
25a.  if (pid != 0) then                                /*process 0 holds global total diff*/
25b.      SEND(mydiff,sizeof(float),0,DIFF);
25c.      RECEIVE(done,sizeof(int),0,DONE);
25d.  else                                              /*pid 0 does this*/
25e.      for i ← 1 to nprocs-1 do                    /*for each other process*/
25f.          RECEIVE(tempdiff,sizeof(float),*,DIFF);
25g.          mydiff += tempdiff;                      /*accumulate into total*/
25h.      endfor
25i.      if (mydiff/(n*n) < TOL) then  done = 1;
25j.      for i ← 1 to nprocs-1 do                    /*for each other process*/
25k.          SEND(done,sizeof(int),i,DONE);
25l.      endfor
25m.  endif
26.  endwhile
27.  end procedure
```

Ocean:

Orchestration in Message Passing

- Message is
 - Sent by calling “Send”
 - Received by calling “Receive”
- Ideally, the sender will wait (block) until the receiver has received the data
 - Thus, communication also establishes synchronization
- There are several “send” and “receive” modes which allow non-blocking behaviour

Orchestration – Message passing

- Blocking Vs non-blocking
 - **Blocking send**
 - Returns when sending buffer is available for reuse
 - MPI_Send
 - **Blocking receive**
 - Waits until data has arrived in receiving buffer
 - MPI_Recv
 - **Non-blocking send and receive**
 - “start sending” and “start receiving”
 - Routine returns immediately, can do other work
 - Should be followed by MPI_Wait
 - $\text{MPI_Send} = \text{MPI_Isend} + \text{MPI_Wait}$
 - $\text{MPI_Recv} = \text{MPI_Ireceive} + \text{MPI_Wait}$

Type of Communication

- Send Modes (MPI Standard, version 3)
 - Standard:
 - **either** sender blocks until matching receive,
 - **or** sender returns right away after copying data into a buffer (safe to re-use original send buffer)
 - Buffered (Bsend):
 - sender returns right away (data copied to new buffer)
 - Synchronous (Ssend):
 - send completes only when matching receive is posted
 - Ready (Rsend):
 - send executes only if the matching receive has already been executed and is ready to receive; otherwise undefined

Type of Communication

- Receive Modes
 - Only one type of receive
 - Blocking: `MPI_Recv()`
 - Non-blocking: `MPI_Irecv()`
 - see also `MPI_Isend()`
 - The “I” means “immediate”, i.e. the system may start to receive data, but call will return before receive is finished allowing concurrent communication

Type of Communication

- Completion
 - MPI_Wait()
 - MPI_Test()
- Probe
 - Check for incoming messages, without actually receiving
 - Determine message size + allocate buffer before receive
 - MPI_Probe() is blocking
 - MPI_IProbe() is non-blocking

Orchestration – Message passing

- **Synchronous vs Asynchronous**
 - Synchronous mode ensures that destination process has started to receive the message
 - Asynchronous doesn't know anything about destination process state

MPI Details

- Read the “MPI Standard”
 - <http://www.mpi-forum.org/docs/>
 - Latest is Version 3.1, June 2015
 - <http://www.mpi-forum.org/docs/mpi-3.1/mpi31-report.pdf>
 - Previous version 3.0, Sept 2012
- Version used on hydra2 is “OpenMPI”
 - Run “`ompi_info`” on command line
 - Installed version is OpenMPI 1.6.5, June 2013
 - (This is not the same as MPI standard version 1.6!!)
 - Implements MPI Standard 2.1

Implementation

WARNING: OpenMP != OpenMPI

- OpenMP is totally different
- OpenMP: shared-memory OpenMPI: message-passing
- OpenMP: you add pragmas to C code (parallelism hints or directives)
- MPI is a general interface for communication between processors that be on entirely different computers (even heterogeneous!)
- **Shared memory programming APIs:**
 - **OpenMP (Open Multi-Processing, we will not use this)**
 - **Pthreads (API for creating and manipulating threads)**
- **Message passing programming APIs:**
 - **OpenMPI (Open Message Passing Interface)**

Assignment 1

- Assignment 1
 - Only do Gaussian Elimination algorithm
 - Do only forward-elimination steps
 - **Parallelize and measure run-time using OpenMPI**
- Read
 - Culler text, Chapter 3
 - Will go over:
 - Message-passing code
 - Shared-memory code

Parallel Programming

Culler, Chapter 3

Culler, Chapter 3

Programming for Performance

- 3.1 Reduce sources of waste
 - Load balancing
 - Reduce communication
 - Reduce serial work (make parallel)
 - Reduce extra work
- Skip 3.2, 3.3, 3.4
- 3.5 Case studies (Ocean, Barnes-Hut, etc)
- Skip 3.6

On your own...

Hennessy & Patterson, Ch. 4

Multiprocessors and Thread-level Parallelism

- Skip 4.1
- 4.2 Symmetric Shared-Memory Architectures (SMPs)
 - Cache coherence
- 4.3 Performance Study of SMPs
- 4.4 Directory-based Coherence
- 4.5 Synchronization (Locks, Barriers)
- Skip 4.6, 4.7, 4.8

Key takeaway

- Two separate issues
 - Cache coherence
 - Memory consistency
- Write a definition for yourself!
 - What is the difference?

